

Generative Adversarial Networks

: From Theory to Practice

Sunmook Choi

Dept. of Mathematics
Korea University

Generative Adversarial Networks

Given a dataset, we want to **generate new data** which seem to be obtained from the dataset, i.e., given p_{data} , we want to learn the **generator's distribution** p_g over data \mathbf{x} s.t. $p_{\text{data}} = p_g$.

How?

- Define a prior $p_z(\cdot)$ on input noise/latent variable \mathbf{z} (e.g., Gaussian), and
- represent a mapping G to a data space as $G(\mathbf{z}; \theta_g)$, which is the **generator**.
- Define the **discriminator** $D(\cdot; \theta_d)$ that outputs a single scalar.
 $D(\mathbf{x})$ represents the probability that \mathbf{x} came from p_{data} rather than p_g .
- Train D to **maximize the probability of assigning the correct label** to both training examples (real, label=1) and samples from G (fake, label=0).
- Simultaneously train G to minimize $\log(1 - D(G(\mathbf{z})))$.

Objective Function

$$\begin{aligned}\min_G \max_D V(D, G) &= \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))] \\ &\approx \frac{1}{m} \sum_{i=1}^m [\log D(\mathbf{x}_i) + \log(1 - D(G(\mathbf{z}_i)))].\end{aligned}$$

Recall: **Binary Cross Entropy Loss** for logistic regression

For binary dataset $\{\mathbf{x}^n, y^n\}_{n=1}^N$ where $y^n \in \{0, 1\}$ with hypothesis model h_{θ} ,

$$BCE(\theta) = \frac{1}{N} \sum_{n=1}^N [-y^n \log(h_{\theta}(\mathbf{x}^n)) - (1 - y^n) \log(1 - h_{\theta}(\mathbf{x}^n))]$$

- For D , $\mathbf{x} \sim p_{\text{data}}(\mathbf{x})$ should have label 1, and $G(\mathbf{z}) \sim p_g(\mathbf{x})$ should have label 0. Hence, maximizing $V(D, G)$ w.r.t. D is equivalent to minimizing BCE.

Objective Function

$$\begin{aligned}\min_G \max_D V(D, G) &= \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))] \\ &\approx \frac{1}{m} \sum_{i=1}^m [\log D(\mathbf{x}_i) + \log(1 - D(G(\mathbf{z}_i)))].\end{aligned}$$

Recall: **Binary Cross Entropy Loss** for logistic regression

For binary dataset $\{\mathbf{x}^n, y^n\}_{n=1}^N$ where $y^n \in \{0, 1\}$ with hypothesis model h_θ ,

$$BCE(\theta) = \frac{1}{N} \sum_{n=1}^N \left[-y^n \log(h_\theta(\mathbf{x}^n)) - (1 - y^n) \log(1 - h_\theta(\mathbf{x}^n)) \right]$$

- Minimizing $V(D, G)$ w.r.t. G may not provide sufficient gradient for G early in learning. Hence, we train G to **minimize $-\log D(G(\mathbf{z}))$** instead of $\log(1 - D(G(\mathbf{z})))$.
- Hence, minimizing $-\log D(G(\mathbf{z}))$ w.r.t. G is equivalent to minimizing BCE.

Theoretical Results

Proposition 1. For G fixed, the optimal discriminator D is $D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})}$.

$$\because V(D, G) = \int_{\mathbf{x}} p_{\text{data}}(\mathbf{x}) \log(D(\mathbf{x})) d\mathbf{x} + \int_{\mathbf{z}} p_z(\mathbf{z}) \log(1 - D(G(\mathbf{z}))) d\mathbf{z} = \int_{\mathbf{x}} p_{\text{data}}(\mathbf{x}) \log(D(\mathbf{x})) + p_g(\mathbf{x}) \log(1 - D(\mathbf{x})) d\mathbf{x}$$

For any $(a, b) \in \mathbb{R}^2 \setminus (0, 0)$, the function $y \mapsto a \log(y) + b \log(1 - y)$ achieves its maximum in $[0, 1]$ at $\frac{a}{a+b}$.

The proposition gives **the analytic form of the optimal discriminator** given a generator G .

With the optimal discriminator, **the objective w.r.t. G** can be reformulated as follows:

$$\begin{aligned} C(G) &= \max_D V(D, G) = V(D_G^*, G) \\ &= \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D_G^*(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z} [\log(1 - D_G^*(G(\mathbf{z})))] \\ &= \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D_G^*(\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim p_g} [\log(1 - D_G^*(\mathbf{x}))] \\ &= \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \left[\log \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})} \right] + \mathbb{E}_{\mathbf{x} \sim p_g} \left[\log \frac{p_g(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})} \right] \end{aligned}$$

Theoretical Results

Theorem 1. The global minimum of $C(G)$ is achieved if and only if $p_g = p_{\text{data}}$.


\because Recall $D_{KL}(p \parallel q) = \mathbb{E}_{x \sim p} \left[\log \frac{p(x)}{q(x)} \right]$ and $D_{JS}(p, q) = \frac{1}{2} D_{KL} \left(p \parallel \frac{p+q}{2} \right) + \frac{1}{2} D_{KL} \left(q \parallel \frac{p+q}{2} \right)$.

$$\begin{aligned} C(G) &= \mathbb{E}_{x \sim p_{\text{data}}} \left[\log \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})} \right] + \mathbb{E}_{x \sim p_g} \left[\log \frac{p_g(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_g(\mathbf{x})} \right] \\ &= D_{KL} \left(p_{\text{data}} \parallel \frac{p_{\text{data}} + p_g}{2} \right) + D_{KL} \left(p_g \parallel \frac{p_{\text{data}} + p_g}{2} \right) - 2 \log 2 \\ &= D_{JS}(p_{\text{data}}, p_g) - \log 4 \end{aligned}$$

Since the Jensen-Shannon divergence is a metric, the only solution is $p_g = p_{\text{data}}$, i.e., the generator G perfectly replicates the data generating process.

Training Algorithm

Algorithm 1 Minibatch stochastic gradient descent training of GAN. The number of steps to apply to the discriminator, k , is a hyperparameter.

- 1: **for** number of epochs **do** **Train discriminator to achieve optimality given a generator.**
- 2: **for** k steps **do** 
- 3: Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_{\mathbf{z}}$.
- 4: Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution p_{data} .
- 5: Update the discriminator by gradient *ascent* method:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log \left(1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

- 6: **end for**
- 7: Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_{\mathbf{z}}$.
- 8: Update the generator by gradient *descent* method:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \left[-\log \left(D(G(\mathbf{z}^{(i)})) \right) \right].$$

- 9: **end for**
-